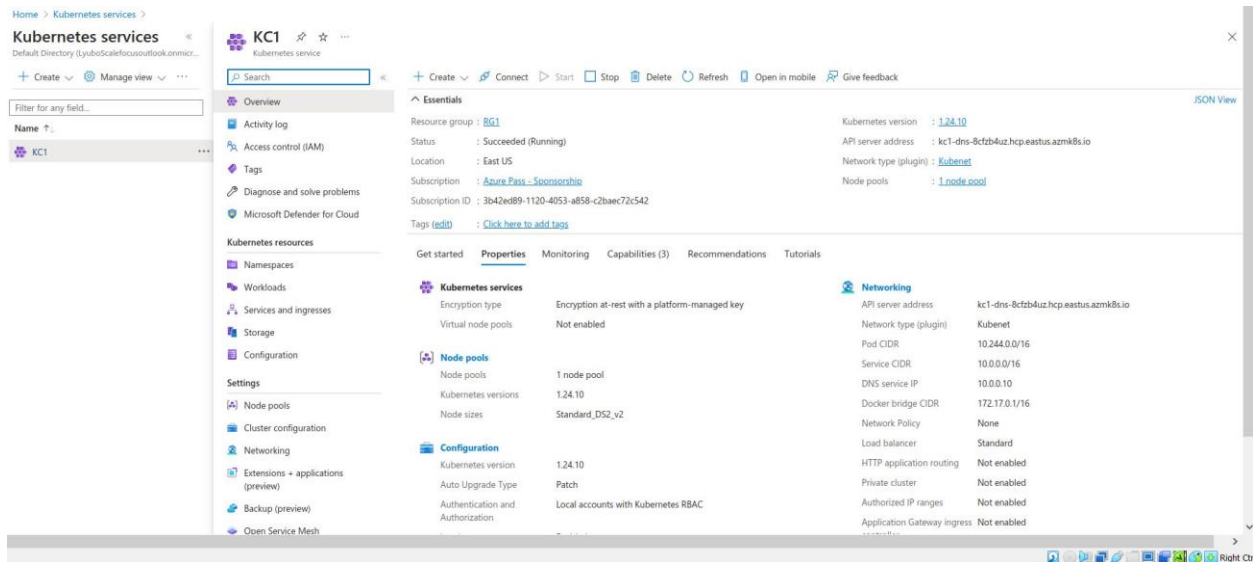


Kubernetes Pods lab

Practice1: Simple pods operations

1. Login to Azure and connect to your AKS cluster.



Now you need to connect to your cluster using command line tooling to interact directly with cluster using kubectl, the command line tool for Kubernetes.

```
lyubomir@mu4a4o:~$ curl -sL https://aka.ms/InstallAzureCLIDeb | sudo bash
[sudo] password for lyubomir:
Hit:1 http://bg.archive.ubuntu.com/ubuntu jammy InRelease
Get:2 http://security.ubuntu.com/ubuntu jammy-security InRelease [110 kB]
Hit:4 https://packages.microsoft.com/repos/azure-cli jammy InRelease
Hit:5 http://bg.archive.ubuntu.com/ubuntu jammy-updates InRelease
Hit:6 http://bg.archive.ubuntu.com/ubuntu jammy-backports InRelease
Get:3 https://packages.cloud.google.com/apt/kubernetes-xenial InRelease [8993 B]
Fetched 119 kB in 1s (115 kB/s)
Reading package lists... Done
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
lib-release is already the newest version (11.1-0ubuntu4).
curl is already the newest version (7.81.0-1ubuntu1.10).
gnupg is already the newest version (2.2.27-3ubuntu2.1).
apt-transport-https is already the newest version (2.4.8).
0 upgraded, 0 newly installed, 0 to remove and 32 not upgraded.
Hit:1 http://bg.archive.ubuntu.com/ubuntu jammy InRelease
Hit:2 http://bg.archive.ubuntu.com/ubuntu jammy-updates InRelease
Hit:3 http://bg.archive.ubuntu.com/ubuntu jammy-backports InRelease
Get:4 http://security.ubuntu.com/ubuntu jammy-security InRelease [110 kB]
Hit:5 https://packages.microsoft.com/repos/azure-cli jammy InRelease
Hit:6 https://packages.cloud.google.com/apt/kubernetes-xenial InRelease
Fetched 110 kB in 1s (113 kB/s)
Reading package lists... Done
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
azure-cli is already the newest version (2.47.0-1-jammy).
0 upgraded, 0 newly installed, 0 to remove and 32 not upgraded.
lyubomir@mu4a4o:~$ az account set --subscription 3b42ed89-1120-4053-a858-c2baec72c542
lyubomir@mu4a4o:~/Desktop$ az aks get-credentials --resource-group RG1 --name KC1
Merged "KC1" as current context in /home/lyubomir/.kube/config
lyubomir@mu4a4o:~/Desktop$ az login
A web browser has been opened at https://login.microsoftonline.com/organizations/oauth2/v2.0/authorize. Please continue the login in the web browser. If no web browser is available or if the web browser
falls to open, use device code flow with 'az login --use-device-code'.
{
  "cloudName": "AzureCloud",
  "homeTenantId": "0dac4ed6-2ff4-485d-a959-43bbf992e367",
  "id": "3b42ed89-1120-4053-a858-c2baec72c542",
  "isDefault": true,
  "managedByTenants": [],
  "name": "Azure Pass - Sponsorship",
  "state": "Enabled",
  "tenantId": "0dac4ed6-2ff4-485d-a959-43bbf992e367",
  "user": {
    "name": "lyubomir@mu4a4o",
    "password": "1234567890",
    "username": "lyubomir@mu4a4o"
  }
}
```

2. Check how many pods run under the default namespace. Run `kubectl get pods`.

```
lyubomir@mu4a4o:~/Desktop$ kubectl get pods
No resources found in default namespace.
```

3. You should not see any pod under the default namespace. Now check all namespaces. Run `kubectl get pods`

`--all-namespaces`.

```
lyubomir@mu4a4o:~/Desktop$ kubectl get pods --all-namespaces
NAMESPACE   NAME                                     READY   STATUS    RESTARTS   AGE
kube-system  ama-logs-rs-57bfc6b798-j52pt           1/1     Running   0           5h11m
kube-system  ama-logs-svhs9                         2/2     Running   0           5h11m
kube-system  azure-ip-masq-agent-svl66             1/1     Running   0           5h11m
kube-system  cloud-node-manager-m7vfz              1/1     Running   0           5h11m
kube-system  coredns-59b6bf8b4f-qf49v              1/1     Running   0           5h11m
kube-system  coredns-59b6bf8b4f-vf95z              1/1     Running   0           5h10m
kube-system  coredns-autoscaler-64b6477b8b-84ftr    1/1     Running   0           5h11m
kube-system  csi-azuredisk-node-t69pj               3/3     Running   0           5h11m
kube-system  csi-azurefile-node-v8v8s               3/3     Running   0           5h11m
kube-system  connectivity-agent-97d4b97f-lfz6s      1/1     Running   0           4h39m
kube-system  connectivity-agent-97d4b97f-plhx4      1/1     Running   0           4h39m
kube-system  kube-proxy-cjhfc                       1/1     Running   0           5h11m
kube-system  metrics-server-7dd74d8758-kjnjg        2/2     Running   0           4h40m
kube-system  metrics-server-7dd74d8758-p5lhm        2/2     Running   0           4h40m
lyubomir@mu4a4o:~/Desktop$
```

4. How many pods do you see? Who deployed these pods? Why are they deployed?

I see fourteen pods. The “kube-system” namespace is a special namespace in Kubernetes that contains system-level components, kubelet and other components that are critical for the operation of the Kubernetes cluster.

5. Now deploy your first pod using the imperative approach. Run `kubectl run nginx --image=nginx`

6. Validate if the pods has been created. What is the status of your pod?

```
lyubomir@mu4a4o:~/Desktop$ kubectl run nginx --image=nginx
pod/nginx created
lyubomir@mu4a4o:~/Desktop$ kubectl get pods
NAME      READY   STATUS    RESTARTS   AGE
nginx     1/1     Running   0           2m46s
lyubomir@mu4a4o:~/Desktop$
```

7. Check the logs coming out of your pod. Run kubectl logs nginx.

```
lyubomir@mu4a4o:~/Desktop$ kubectl logs nginx
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf
10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/30-tune-worker-processes.sh
/docker-entrypoint.sh: Configuration complete; ready for start up
2023/04/15 15:42:40 [notice] 1#1: using the "epoll" event method
2023/04/15 15:42:40 [notice] 1#1: nginx/1.23.4
2023/04/15 15:42:40 [notice] 1#1: built by gcc 10.2.1 20210110 (Debian 10.2.1-6)
2023/04/15 15:42:40 [notice] 1#1: OS: Linux 5.4.0-1105-azure
2023/04/15 15:42:40 [notice] 1#1: getrlimit(RLIMIT_NOFILE): 1048576:1048576
2023/04/15 15:42:40 [notice] 1#1: start worker processes
2023/04/15 15:42:40 [notice] 1#1: start worker process 29
2023/04/15 15:42:40 [notice] 1#1: start worker process 30
lyubomir@mu4a4o:~/Desktop$
```

8. Run following command to check current resource consumption of your pod: kubectl top pod nginx.

9. Check on which Node your pods has been scheduled. Run kubectl get pods -o wide.

10. Try to find the same information but this time running kubectl describe pod nginx.

```
lyubomir@mu4a4o:~/Desktop$ kubectl get pods -o wide
NAME      READY   STATUS    RESTARTS   AGE   IP           NODE                                     NOMINATED NODE   READINESS GATES
nginx     1/1     Running   0           9m1s  10.244.0.18  aks-agentpool-35267590-vmss000002    <none>           <none>
lyubomir@mu4a4o:~/Desktop$ kubectl describe pod nginx
Name:      nginx
Namespace: default
Priority:   0
Service Account: default
Node:      aks-agentpool-35267590-vmss000002/10.224.0.6
Start Time: Sat, 15 Apr 2023 18:42:37 +0300
Labels:    run=nginx
Annotations: <none>
Status:    Running
IP:        10.244.0.18
IPs:
  IP: 10.244.0.18
Containers:
  nginx:
    Container ID:   containerd://a51f0c006a01300009037b3744e51cd101653f08c8839efc98d41784206e0d05
    Image:          nginx
    Image ID:       docker.io/library/nginx@sha256:63b44e8ddb83d5dd8020327c1f40436e37a6fffd3ef2498a6204df23be6e7e94
    Port:           <none>
    Host Port:      <none>
    State:          Running
      Started:      Sat, 15 Apr 2023 18:42:40 +0300
      Ready:        True
      Restart Count: 0
    Environment:    <none>
```



```

Image:      nginx
Image ID:   docker.io/library/nginx@sha256:63b44e8ddb83d5dd8020327c1f40436e37a6fffd3ef2498a6204df23be67e94
Port:       <none>
Host Port:  <none>
State:      Running
  Started:   Sat, 15 Apr 2023 18:42:40 +0300
Ready:      True
Restart Count: 0
Environment: <none>
Mounts:
  /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-99lvv (ro)
Conditions:
  Type             Status
  Initialized       True
  Ready             True
  ContainersReady   True
  PodScheduled      True
Volumes:
  kube-api-access-99lvv:
    Type:      Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName: kube-root-ca.crt
    ConfigMapOptional: <nil>
    DownwardAPI: true
QoS Class:      BestEffort
Node-Selectors: <none>
Tolerations:    node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
                 node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type     Reason      Age   From          Message
  ----     -
  Normal   Scheduled   10m   default-scheduler   Successfully assigned default/nginx to aks-agentpool-35267590-vmss000002
  Normal   Pulling     10m   kubelet          Pulling image "nginx"
  Normal   Pulled      10m   kubelet          Successfully pulled image "nginx" in 2.606398353s
  Normal   Created     10m   kubelet          Created container nginx
  Normal   Started     10m   kubelet          Started container nginx
lyubomir@mu4a4o:~/Desktop$

```

11. Delete your pod using kubectl delete pod nginx.

```

lyubomir@mu4a4o:~/Desktop$ kubectl delete pod nginx
pod "nginx" deleted

```

12. Let's find the image used on one of the coredns pods under the kube-system namespace.

```

lyubomir@mu4a4o:~/Desktop$ kubectl get pods --all-namespaces

```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	ama-logs-rs-57bfc6b798-j52pt	1/1	Running	0	5h11m
kube-system	ama-logs-svhs9	2/2	Running	0	5h11m
kube-system	azure-ip-masq-agent-svl66	1/1	Running	0	5h11m
kube-system	cloud-node-manager-m7vfz	1/1	Running	0	5h11m
kube-system	coredns-59b6bf8b4f-qf49v	1/1	Running	0	5h11m
kube-system	coredns-59b6bf8b4f-vf95z	1/1	Running	0	5h10m
kube-system	<u>coredns-autoscaler-64b6477b8b-84ftr</u>	1/1	Running	0	5h11m
kube-system	csi-azuredisk-node-t69pj	3/3	Running	0	5h11m
kube-system	csi-azurefile-node-v8v8s	3/3	Running	0	5h11m
kube-system	konnnectivity-agent-97d4b97f-lfz6s	1/1	Running	0	4h39m
kube-system	konnnectivity-agent-97d4b97f-plhx4	1/1	Running	0	4h39m
kube-system	kube-proxy-cjhfc	1/1	Running	0	5h11m
kube-system	metrics-server-7dd74d8758-kjnjg	2/2	Running	0	4h40m
kube-system	metrics-server-7dd74d8758-p5lhm	2/2	Running	0	4h40m

```

lyubomir@mu4a4o:~/Desktop$

```

13. Once again list all pods under all namespaces.

14. Note one of the coredns pods. Now run kubectl describe pod <coredns-name> -n kube-system. Replace the <coredns-name> place holder with noted name.

15. Inspect the output and locate the image information.

```

lyubomir@mu4a4o: ~/Desktop$ kubectl describe pod coredns-59b6bf8b4f-vf95z -n kube-system
Name: coredns-59b6bf8b4f-vf95z
Namespace: kube-system
Priority: 2000001000
Priority Class Name: system-node-critical
Service Account: coredns
Node: aks-agentpool-35267590-vmss000002/10.224.0.6
Start Time: Sat, 15 Apr 2023 13:28:14 +0300
Labels: k8s-app=kube-dns
        kubernetes.io/cluster-service=true
        pod-template-hash=59b6bf8b4f
        version=v20
Annotations: prometheus.io/port: 9153
Status: Running
IP: 10.244.0.10
IPs:
  IP: 10.244.0.10
Controlled By: ReplicaSet/coredns-59b6bf8b4f
Containers:
  coredns:
    Container ID: containerd://cb88f64ce1a565d495fa29b1ba46e65be55f4335c4fb474d73fb4815cd480704
    Image: mcr.microsoft.com/oss/kubernetes/coredns:v1.9.3
    Image ID: sha256:c38f956b642366c8eeb0babfda6b0bb2aa92f27a968589804cadb445f6df72d6
    Ports: 53/UDP, 53/TCP, 9153/TCP
    Host Ports: 0/UDP, 0/TCP, 0/TCP
    Args:
      -conf
      /etc/coredns/Corefile
    State: Running
      Started: Sat, 15 Apr 2023 13:28:14 +0300
    Ready: True
    Restart Count: 0
    Limits:
      cpu: 3
      memory: 500Mi
    Requests:
      cpu: 100m
      memory: 70Mi
    Liveness: http-get http://:8080/health delay=60s timeout=5s period=10s #success=1 #failure=5
    Readiness: http-get http://:8181/ready delay=0s timeout=1s period=10s #success=1 #failure=3
    Environment:
      KUBERNETES_PORT_443_TCP_ADDR: kci-dns-8cfzb4uz.hcp.eastus.azmk8s.io
      KUBERNETES_PORT_443_TCP: tcp://kci-dns-8cfzb4uz.hcp.eastus.azmk8s.io:443
      KUBERNETES_SERVICE_HOST: kci-dns-8cfzb4uz.hcp.eastus.azmk8s.io
    Mounts:
      /etc/coredns from config-volume (ro)
    Mounts:
      /etc/coredns from config-volume (ro)
      /etc/coredns/custom from custom-config-volume (ro)
      /tmp from tmp (rw)
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-g48td (ro)
    Conditions:
      Type Status
      Initialized True
      Ready True
      ContainersReady True
      PodScheduled True
    Volumes:
      config-volume:
        Type: ConfigMap (a volume populated by a ConfigMap)
        Name: coredns
        Optional: false
      custom-config-volume:
        Type: ConfigMap (a volume populated by a ConfigMap)
        Name: coredns-custom
        Optional: true
      tmp:
        Type: EmptyDir (a temporary directory that shares a pod's lifetime)
        Medium:
        SizeLimit: <unset>
      kube-api-access-g48td:
        Type: Projected (a volume that contains injected data from multiple sources)
        TokenExpirationSeconds: 3607
        ConfigMapName: kube-root-ca.crt
        ConfigMapOptional: <nil>
        DownwardAPI: true
    QoS Class: Burstable
    Node-Selectors: <none>
    Tolerations: CriticalAddonsOnly op=Exists
      node-role.kubernetes.io/master:NoSchedule
      node.kubernetes.io/memory-pressure:NoSchedule op=Exists
      node.kubernetes.io/not-ready:NoExecute op=Exists for 30s
      node.kubernetes.io/unreachable:NoExecute op=Exists for 30s
    Events:
lyubomir@mu4a4o: ~/Desktop$

```

16. Now let us check the logs of the metrics-server pod. Run the same command as in step 7 but don't forget to add the namespace in which this pod is created.

```
lyubonir@mu4a4o:~/Desktop$ kubectl get pods --all-namespaces
NAMESPACE      NAME                                     READY   STATUS    RESTARTS   AGE
kube-system    ana-logs-rs-57bfceb798-j52pt           1/1     Running   0           20h
kube-system    ana-logs-svhs9                          2/2     Running   0           20h
kube-system    azure-ip-masq-agent-svl66              1/1     Running   0           20h
kube-system    cloud-node-manager-n7vfz               1/1     Running   0           20h
kube-system    coredns-59bbbf8b4f-qf49v               1/1     Running   0           20h
kube-system    coredns-59bbbf8b4f-vf95z               1/1     Running   0           20h
kube-system    coredns-autoscaler-64b6477b8b-84ftr    1/1     Running   0           20h
kube-system    csi-azuredisk-node-t69pj                3/3     Running   0           20h
kube-system    csi-azurefile-node-v6d8s                3/3     Running   0           20h
kube-system    connectivity-agent-97d4b97f-lfz6s       1/1     Running   0           19h
kube-system    connectivity-agent-97d4b97f-plhx4       1/1     Running   0           19h
kube-system    kube-proxy-cjhfc                        1/1     Running   0           20h
kube-system    metrics-server-7dd74d8758-kjnlg         2/2     Running   0           19h
kube-system    metrics-server-7dd74d8758-pslhm         2/2     Running   0           19h
lyubonir@mu4a4o:~/Desktop$ kubectl logs metrics-server-7dd74d8758-pslhm -n kube-system
Defaulted container "metrics-server-vpa" out of: metrics-server-vpa, metrics-server
ERROR: logging before flag.Parse: I0415 10:58:24.352244      1 pod_nanny.go:68] Invoked by [/pod_nanny --config-dir=/etc/config --cpu=44m --extra-cpu=0.5m --memory=51Ml --extra-memory=4Ml --poll-per
iod=180000 --thresholds --deployment-metrics-server --container-metrics-server]
ERROR: logging before flag.Parse: I0415 10:58:24.352295      1 pod_nanny.go:69] Version: 1.8.14
ERROR: logging before flag.Parse: I0415 10:58:24.352308      1 pod_nanny.go:85] Watching namespace: kube-system, pod: metrics-server-7dd74d8758-pslhm, container: metrics-server.
ERROR: logging before flag.Parse: I0415 10:58:24.352313      1 pod_nanny.go:86] storage: MISSING, extra_storage: 0G
ERROR: logging before flag.Parse: I0415 10:58:24.353318      1 pod_nanny.go:189] Failed to read data from config file "/etc/config/NannyConfiguration": open /etc/config/NannyConfiguration: no such f
ile or directory, using default parameters
ERROR: logging before flag.Parse: I0415 10:58:24.353350      1 pod_nanny.go:116] cpu: 44m, extra_cpu: 0.5m, memory: 51Ml, extra_memory: 4Ml
ERROR: logging before flag.Parse: I0415 10:58:24.353360      1 pod_nanny.go:145] Resources: [(Base:(t:(value:44 scale:-3) d:(Dec:cnl) s:44m Format:DecimalSI) ExtraPerNode:(t:(value:5 scale:-4) d:(
Dec:cnl) s: Format:DecimalSI) Name:cpu) (Base:(t:(value:53477376 scale:0) d:(Dec:cnl) s:51Ml Format:BinarySI) ExtraPerNode:(t:(value:4194304 scale:0) d:(Dec:cnl) s:4Ml Format:BinarySI) Name:mem
ory)]
lyubonir@mu4a4o:~/Desktop$
```

Practice2: Working with pod manifest files

Note: Try not to do a copy/paste on commands requests unless you are instructed to do so.

Copy/paste will not help you to learn Kubernetes!

1. Now it is time to deploy pod using manifest file (declarative approach). Copy the following code block on your local computer in a file called redis.yaml:

apiVersion: v11

kind: pod

metadata:

name: static-web

labels:

role: myrole

specs:

containers:

- name: redis

image: redis123

2. Try to deploy the pod defined in redis.yaml. Run `kubectl create -f redis.yaml`.

3. You will receive errors on your screen. Your next task will be to correct the syntax of the code you just

copied. You can use the online Kubernetes documentation or you can search the internet in general.

```
lyubomir@mu4a4o:~/Desktop$ kubectl create -f redis.yaml
Unable to connect to the server: dial tcp: lookup kc1-dns-8cfzb4uz.hcp.eastus.azmk8s.io on 127.0.0.53:53: no such host
lyubomir@mu4a4o:~/Desktop$
```

4. When you solve all the syntax errors your pod should be deployed but is it running? What is the status of your pod?

```
lyubomir@mu4a4o:~/Desktop$ kubectl create -f redis.yaml
pod/static-web created
lyubomir@mu4a4o:~/Desktop$
```

```
lyubomir@mu4a4o:~/Desktop$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
static-web    1/1     Running   0           2m21s
lyubomir@mu4a4o:~/Desktop$
```

5. Check the events associated with this pod. Run the `kubectl describe pod static-web` command. What are the events showing? Why your pod is not running?

6. Find the correct image (check the Docker hub page) and correct it in the manifest.

7. Locate the image information and put the correct image name. Redeploy the pod (first run `kubectl delete pod static-web` to delete the pod, then run `kubectl create` once again).

8. Check the status of your pod. It should be running now.

```
lyubomir@mu4a4o:~/Desktop$ kubectl describe pod static-web
Name:          static-web
Namespace:     default
Priority:       0
Service Account: default
Node:          aks-agentpool-28149888-vmss000001/10.224.0.4
Start Time:    Sun, 16 Apr 2023 22:04:31 +0300
Labels:        name=myrole
Annotations:    <none>
Status:        Running
IP:            10.244.2.11
IPs:           IP: 10.244.2.11
Containers:
  redis:
    Container ID:  containerd://5b149028ce96cec3c37c61aeea02c74de70aa31873558c3848f8fcb80ee58e2f
    Image:         redis
    Image ID:      docker.io/library/redis@sha256:92b8b307ee28ed74da17578064c73307ad41e43f422f0b7e4e91498b406c59e3
    Port:         <none>
    Host Port:    <none>
    State:        Running
      Started:    Sun, 16 Apr 2023 22:04:35 +0300
    Ready:        True
    Restart Count: 0
    Environment:  <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-88c9r (ro)
Conditions:
  Type              Status
  Initialized       True
  Ready             True
  ContainersReady   True
  PodScheduled      True
Volumes:
```



```

Volumes:
  kube-api-access-88c9r:
    Type: Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName: kube-root-ca.crt
    ConfigMapOptional: <nil>
    DownwardAPI: true
QoS Class: BestEffort
Node-Selectors: <none>
Tolerations: node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
              node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type     Reason      Age   From          Message
  ----     -
  Normal   Scheduled   66s   default-scheduler   Successfully assigned default/static-web to aks-agentpool-28149888-vmss000001
  Normal   Pulling     65s   kubelet          Pulling image "redis"
  Normal   Pulled      62s   kubelet          Successfully pulled image "redis" in 2.436374163s
  Normal   Created     62s   kubelet          Created container redis
  Normal   Started     62s   kubelet          Started container redis

```

9. Now you can delete the pod. Try to delete it using the `kubectl delete -f redis.yaml`.

```

lyubomir@mu4a4o:~/Desktop$ kubectl delete -f redis.yaml
pod "static-web" deleted
lyubomir@mu4a4o:~/Desktop$

```

10. Your next task is to create and test nginx pod definition. Your definition should use the nginx official image, should use label named `app` with value `frontend` and should publish port 80. Make sure you complete this task because we will use this template in our next Labs. Your nginx pod should be running without any issues.

```

lyubomir@mu4a4o:~/Desktop$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
nginredis     2/2     Running   0           23m
nginx-pod     1/1     Running   0           81m
lyubomir@mu4a4o:~/Desktop$

```

11. Final task of this practice will be to define pod definition with following details:

- Image=memcached
- Port= 11211
- Label app=web
- CPU request=0.35 cores
- RAM request=0.15 GB
- CPU limit=0.5 cores
- Ram limit=0.25 GB
- Restart policy=Never

12. Don't forget to try your pod definition.


```
lyubomir@mu4a4o:~/Desktop$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
memcached     1/1     Running   0           9s
nginredis     2/2     Running   0          33m
nginx-pod     1/1     Running   0          90m
lyubomir@mu4a4o:~/Desktop$
```

Practice3: Multi-container pods

1. Once finished you can try to create multi-container pod definition. Your multi-container pod should use redis and nginx containers with port 6379 and 80 published respectively. Label name should be app with value web.
2. Note that in reality there is no sense to put the redis and nginx under the same pod but it can be done for the purpose of learning.
3. Deploy your multi-container pod. It should have running status. What is written under Ready column when you kubectl get the pods? Why your pod displays different values for ready?

```
lyubomir@mu4a4o:~/Desktop$ kubectl apply -f multicontainer.yaml
pod/nginredis unchanged
lyubomir@mu4a4o:~/Desktop$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
memcached     1/1     Running   0           3m1s
nginredis     2/2     Running   0           35m
nginx-pod     1/1     Running   0           93m
```

4. Kubectl describe you new pod, and locate the containers section. How many containers are listed?
5. Delete all the pods under the default namespace.

```
lyubomir@mu4a4o:~/Desktop$ kubectl get pods -namespace
No resources found in amespaces namespace.
```

6. Don't delete any of the manifest files you have created so far.

Practice4: Probes

1. First we will create and test liveness probe with exec test. Create a file named probes_exec.yaml with following content:

apiVersion: v1

kind: Pod

```

metadata:

labels:

test: liveness

name: liveness-exec

spec:

containers:

- name: liveness

image: k8s.gcr.io/busybox

args:

- /bin/sh

- -c

- touch /tmp/healthy; sleep 30; rm -rf /tmp/healthy; sleep 600

livenessProbe:

exec:

command:

- cat

- /tmp/healthy

initialDelaySeconds: 5

periodSeconds: 5

```

2. Examine the containers args commands especially the line that start with touch. This bash pipeline will help us to test the liveness probes.

3. Run `kubectl create -f probes_exec.yaml`.

```

lyubomir@mu4a4o:~/Desktop$ kubectl create -f probes_exec.yaml
pod/liveness-exec created
lyubomir@mu4a4o:~/Desktop$

```

4. Run `kubectl describe pod liveness-exec` immediately after you deploy the pod. The output should indicate that no liveness probes have failed yet.

```

Events:
  Type    Reason      Age    From              Message
  ----    -
  Normal  Scheduled   104s   default-scheduler Successfully assigned default/liveness-exec to aks-agentpool-28149888-vmss000001
  Normal  Pulled      103s   kubelet           Successfully pulled image "k8s.gcr.io/busybox" in 521.527522ms
  Warning  Unhealthy   59s (x3 over 69s)  kubelet           Liveness probe failed: cat: can't open '/tmp/healthy': No such file or directory
  Normal  Killing     59s   kubelet           Container liveness failed liveness probe, will be restarted
  Normal  Pulling     29s (x2 over 104s)  kubelet           Pulling image "k8s.gcr.io/busybox"
  Normal  Created     29s (x2 over 103s)  kubelet           Created container liveness
  Normal  Started     29s (x2 over 103s)  kubelet           Started container liveness
  Normal  Pulled      29s   kubelet           Successfully pulled image "k8s.gcr.io/busybox" in 297.268054ms

```

5. After 35 seconds, view the Pod events again. Run `kubectl describe pod liveness-exec`.
6. At the bottom of the output, there should be a messages indicating that the liveness probes have failed, and the containers have been killed and recreated.

Events:				
Type	Reason	Age	From	Message
Normal	Scheduled	66s	default-scheduler	Successfully assigned default/liveness-exec to aks-agentpool-28149888-vmss000001
Normal	Pulling	65s	kubelet	Pulling image "k8s.gcr.io/busybox"
Normal	Pulled	64s	kubelet	Successfully pulled image "k8s.gcr.io/busybox" in 521.527522ms
Normal	Created	64s	kubelet	Created container liveness
Normal	Started	64s	kubelet	Started container liveness
Warning	Unhealthy	20s (x3 over 30s)	kubelet	Liveness probe failed: cat: can't open '/tmp/healthy': No such file or directory
Normal	Killing	20s	kubelet	Container liveness failed liveness probe, will be restarted

7. Wait another 30 seconds, and verify that the container has been restarted. Run `kubectl get pod livenessexec`.
8. The output should show that `RESTARTS` has been incremented.

9. We will continue with HTTP probe. Create file named `probes_http.yaml` with following content:

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
labels:
```

```
test: liveness
```

```
name: liveness-http
```

```
spec:
```

```
containers:
```

```
- name: liveness
```

```
image: k8s.gcr.io/liveness
```

```
args:
```

```
- /server
```

```
livenessProbe:
```

```
httpGet:
```

```
path: /healthz
```

```
port: 8080
```

```
httpHeaders:
```

```
- name: Custom-Header
```

```
value: Awesome
```

initialDelaySeconds: 3

periodSeconds: 3

10. Just for your info, /healthz handler has following function implemented:

```
http.HandleFunc("/healthz", func(w http.ResponseWriter, r *http.Request) {  
    duration := time.Now().Sub(started)  
    if duration.Seconds() > 10 {  
        w.WriteHeader(500)  
        w.Write([]byte(fmt.Sprintf("error: %v", duration.Seconds())))  
    } else {  
        w.WriteHeader(200)  
        w.Write([]byte("ok"))  
    }  
})
```

11. For the first 10 seconds that the container is alive, the /healthz handler returns a status of 200. After that, the handler returns a status of 500.

12. Run `kubectl create -f probes_http.yaml`.

```
lyubomir@mu4a4o:~/Desktop$ kubectl create -f probes_http.yaml  
pod/liveness-http created
```

13. Immediately run (you only have 10 secs to run this command) `kubectl describe pod liveness-http`.

14. Your pod should be live and running.

15. After 10 seconds, view Pod events to verify that liveness probes have failed and the container has been restarted. Run again `kubectl describe pod liveness-http`.

16. You should see the same output as in step 7. Kubelet will reboot the container.

Events:	Reason	Age	From	Message
Normal	Scheduled	114s	default-scheduler	Successfully assigned default/liveness-http to aks-agentpool-28149888-vmss000001
Normal	Pulled	113s	kubelet	Successfully pulled image "k8s.gcr.io/liveness" in 589.162723ms
Normal	Pulled	96s	kubelet	Successfully pulled image "k8s.gcr.io/liveness" in 336.498626ms
Normal	Created	78s (x3 over 113s)	kubelet	Created container liveness
Normal	Started	78s (x3 over 113s)	kubelet	Started container liveness
Normal	Pulled	78s	kubelet	Successfully pulled image "k8s.gcr.io/liveness" in 297.022979ms
Normal	Pulling	60s (x4 over 114s)	kubelet	Pulling image "k8s.gcr.io/liveness"
Warning	Unhealthy	60s (x9 over 102s)	kubelet	Liveness probe failed: HTTP probe failed with statuscode: 500
Normal	Killing	60s (x3 over 96s)	kubelet	Container liveness failed liveness probe, will be restarted

17. We continue with TCP probes. Create file named `probes_tcp.yaml` with following content:

apiVersion: v1

kind: Pod

metadata:

name: liveness-tcp

labels:

app: goproxy

spec:

containers:

- name: goproxy

image: k8s.gcr.io/goproxy:0.1

ports:

- containerPort: 8080

livenessProbe:

tcpSocket:

port: 9999 #8080 is valid port

initialDelaySeconds: 15

periodSeconds: 20

18. Run `kubectl create -f probes_tcp.yaml`.

```
lyubomir@mu4a4o:~/Desktop$ kubectl create -f probes_tcp.yaml
pod/liveness-tcp created
```

19. Immediately run (you only have 10 secs to run this command) `kubectl describe pod liveness-tcp`.

20. Your pod should be live and running.

```
Events:
  Type     Reason      Age   From          Message
  ----     -
  Normal   Scheduled   5s    default-scheduler   Successfully assigned default/liveness-tcp to aks-agentpool-28149888-vmss000001
  Normal   Pulling     5s    kubelet        Pulling image "k8s.gcr.io/goproxy:0.1"
  Normal   Pulled      4s    kubelet        Successfully pulled image "k8s.gcr.io/goproxy:0.1" in 797.555461ms
  Normal   Created     4s    kubelet        Created container goproxy
  Normal   Started     4s    kubelet        Started container goproxy
```

21. After 10 seconds, view Pod events to verify that liveness probes have failed and the container has been restarted. Run again `kubectl describe pod liveness-tcp`.

```
Events:
  Type     Reason      Age   From          Message
  ----     -
  Normal   Scheduled   26s   default-scheduler   Successfully assigned default/liveness-tcp to aks-agentpool-28149888-vmss000001
  Normal   Pulling     26s   kubelet        Pulling image "k8s.gcr.io/goproxy:0.1"
  Normal   Pulled      25s   kubelet        Successfully pulled image "k8s.gcr.io/goproxy:0.1" in 797.555461ms
  Normal   Created     25s   kubelet        Created container goproxy
  Normal   Started     25s   kubelet        Started container goproxy
  Warning  Unhealthy   6s    kubelet        Liveness probe failed: dial tcp 10.244.2.25:9999: connect: connection refused
```

22. You should see the same output as in step 7 and 16. Kubelet will reboot the container.

23. Our last job will be to define one readiness probe using HTTP test.

24. Create file named readiness_http.yaml with following content:

```
apiVersion: v1
kind: Pod
metadata:
  name: readiness-http
  labels:
    app: test
spec:
  containers:
    - name: nginx
      image: nginx
      ports:
        - containerPort: 80
      readinessProbe:
        initialDelaySeconds: 1
        periodSeconds: 2
        timeoutSeconds: 1
        successThreshold: 1
        failureThreshold: 1
        httpGet:
          host:
          scheme: HTTP
          path: /
          httpHeaders:
            - name: Host
              value: myapplication1.com
          port: 80
```

25. Run `kubectl create -f readiness_http.yaml`.

26. Run `kubectl get pods -A` to see the status of your pod.
27. Pods and their status and ready states will be displayed; our pod should be in running state.
28. Run `kubectl describe pod readiness-http`. Examine the events for this pod. Everything should be OK.
29. Now delete the pod and edit the `readiness_http.yaml` so that the port parameter has 81 value.
30. Run again `kubectl create -f readiness_http.yaml`.
31. Run `kubectl get pods -A` to see the status of your pod. You should see that the pod is running but it is not in ready state.

```
lyubomir@mu4a4o:~/Desktop$ kubectl create -f readiness_http.yaml
pod/readiness-http created
lyubomir@mu4a4o:~/Desktop$ kubectl get pods -A
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
default	liveness-exec	0/1	CrashLoopBackOff	21 (24s ago)	64m
default	liveness-http	1/1	Running	25 (14s ago)	58m
default	liveness-tcp	0/1	CrashLoopBackOff	9 (40s ago)	20m
default	memcached	1/1	Running	0	72m
default	nginxredis	2/2	Running	0	105m
default	nginx-pod	1/1	Running	0	162m
default	readiness-http	0/1	Running	0	8s
kube-system	ama-logs-b4pax	2/2	Running	0	4h39m
kube-system	ama-logs-rs-b9b64776-sv8j8	1/1	Running	0	4h39m
kube-system	azure-ip-masq-agent-nl92v	1/1	Running	0	4h39m
kube-system	cloud-node-manager-2d665	1/1	Running	0	4h39m
kube-system	coredns-59b6bf8b4f-j6tcz	1/1	Running	0	4h38m
kube-system	coredns-59b6bf8b4f-mbjzf	1/1	Running	0	4h39m
kube-system	coredns-autoscaler-64b6477b8b-z96nt	1/1	Running	0	4h39m
kube-system	csi-azuredisk-node-n9v8g	3/3	Running	0	4h39m
kube-system	csi-azurefile-node-5hhkv	3/3	Running	0	4h39m
kube-system	konnectivity-agent-656f67bbdd-hv8nc	1/1	Running	0	3h44m
kube-system	konnectivity-agent-656f67bbdd-zcfqx	1/1	Running	0	3h44m
kube-system	kube-proxy-d2nq6	1/1	Running	0	4h39m
kube-system	metrics-server-7dd74d8758-brgpp	2/2	Running	0	4h8m
kube-system	metrics-server-7dd74d8758-gnwzc	2/2	Running	0	4h8m

```
lyubomir@mu4a4o:~/Desktop$
```

32. Describe the pod. Run `kubectl describe pod readiness-http`.
33. From the events we can see that readiness probe failed due to the connection being refused therefore pod will not receive any traffic.
34. Delete all pods under the default namespace.
35. Don't delete any manifest files created so far.